

INTRO TO DATA SCIENCE RECOMMENDATION ENGINES

I. DATA TYPES

II. CONTENT-BASED FILTERING

III. COLLABORATIVE FILTERING

EXERCISE:

IV. RECOMMENDING WITH PYTHON

V. THE NETFLIX PRIZE

A recommendation system aims to match users to products/items/brand/etc that they likely haven't experienced yet and/or predict a user's preference based on past observations.

A recommendation system aims to match users to products/items/brand/etc that they likely haven't experienced yet and/or predict a user's preference based on past observations.

A **ranking** or **prediction** is produced by analyzing other user/item ratings (and sometimes item characteristics) to provide personalized recommendations to users.

INTRO TO DATA SCIENCE

I. TYPES OF DATA

EXAMPLES – TYPES OF DATA

THE KIND OF RECOMMENDATIONS YOU CAN GIVE, ARE DEPENDENT ON THE DATA YOU HAVE.

Inspired by Your Shopping Trends



WE NEED DATA TO RECOMMEND.

- Preferences
- Ratings
- Item meta-data
- User Behavior



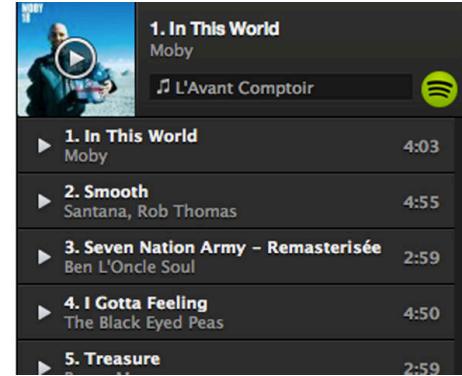
EXAMPLES – TYPES OF DATA

Excellent	
Very Good	
Good	
Fair	
Poor	
No rating submitted	N/A

User Log

Name	Email	Role	Login	IP Address	Duration
John Bob	johnbob@chide.it	Member	Apr 3, 4:01 PM	127.0.0.1	0:00:05
John Bob	johnbob@chide.it	Member	Apr 3, 2:58 PM	127.0.0.1	0:00:42
Jane Doe	janedoe@chide.it	Member	Apr 17, 3:02 PM	127.0.0.1	0:00:04
Jane Doe	janedoe@chide.it	Member	Apr 17, 2:54 PM	127.0.0.1	0:02:30
Frank Storm	frankstorm@chide.it	Admin	Apr 17, 3:02 PM	127.0.0.1	0:00:04
Frank Storm	frankstorm@chide.it	Admin	Apr 17, 3:00 PM	127.0.0.1	0:01:28
Adam Klockars	adam@chide.it	Admin	Apr 17, 3:02 PM	127.0.0.1	Active
Adam Klockars	adam@chide.it	Admin	Apr 17, 3:01 PM	127.0.0.1	0:00:03
Adam Klockars	adam@chide.it	Admin	Apr 17, 3:01 PM	127.0.0.1	0:00:06
Adam Klockars	adam@chide.it	Admin	Apr 17, 2:59 PM	127.0.0.1	0:00:48

1 - 10 of 22 > [10 20 50]



Ratings
Upvotes / Downvotes
Weighted Scale
Grades
Relevance Feedback

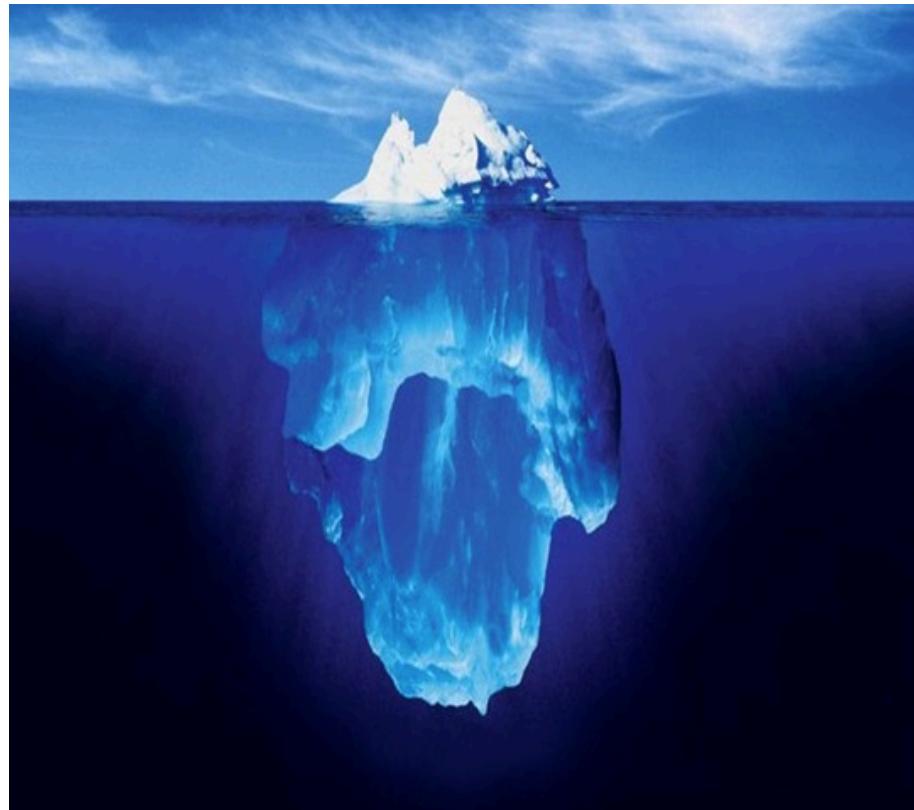
Access Logs
Session Lengths
Time spent on a page
Clicks / Non-Clicks
Purchase History
Product Descriptions

Listening History
Playlist Creates
Follows / Unfriend
Impressions
Email Reads / Impressions

EXAMPLES – TYPES OF DATA

9

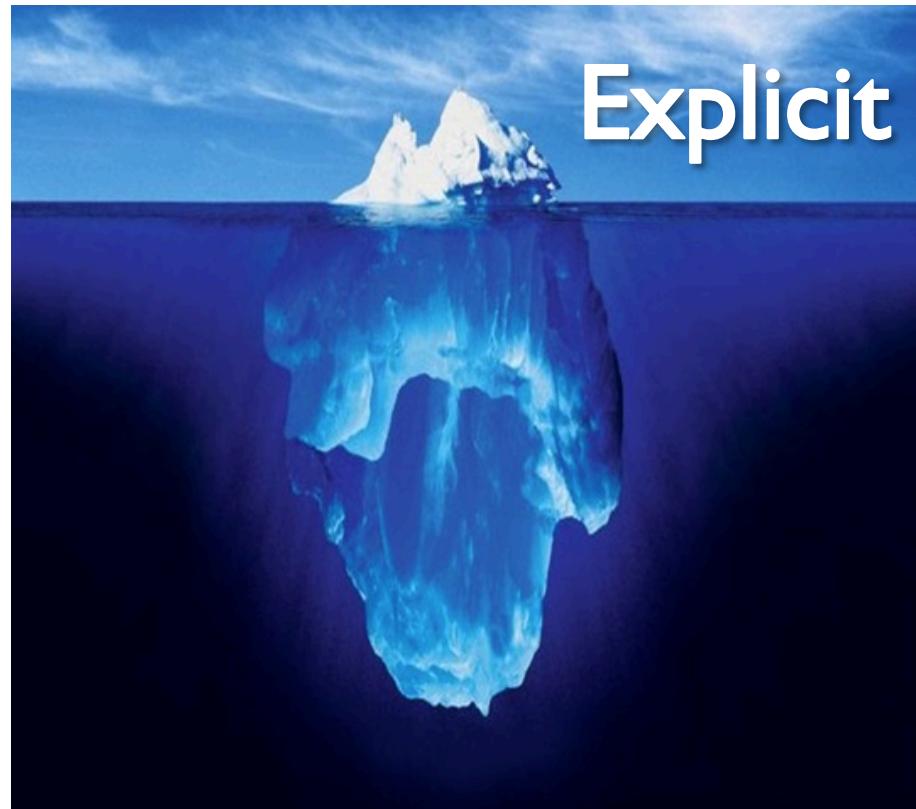
Recommenders need feedback to be useful.



Recommenders need feedback to be useful.

Explicit

- Explicitly given
- Pro-actively acquired
- Expensive to collect



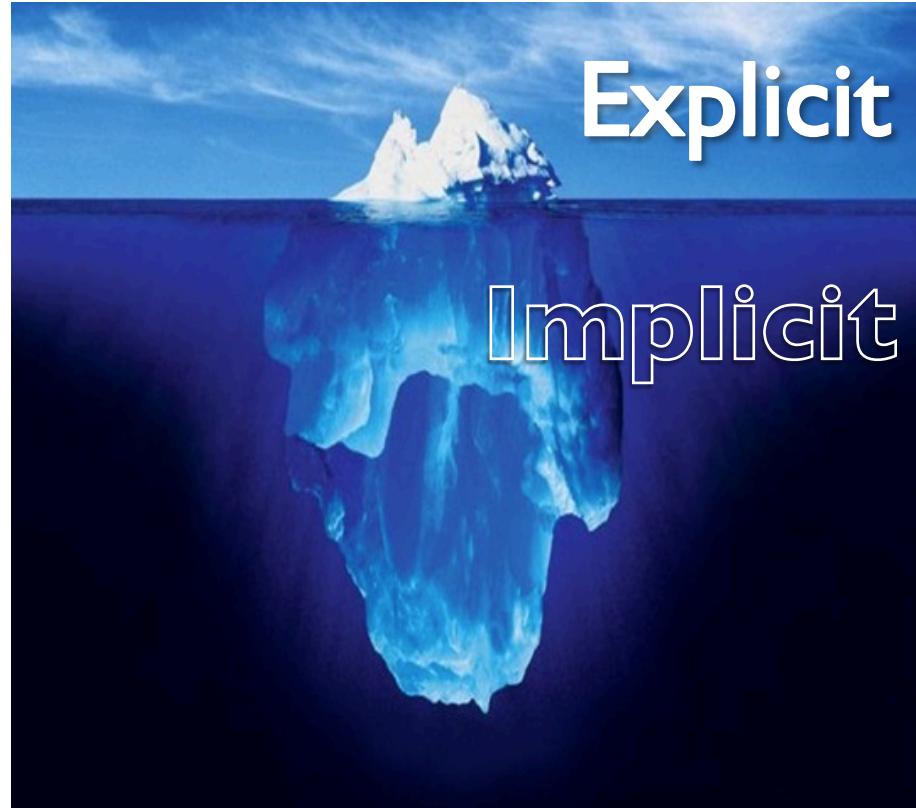
Recommenders need feedback to be useful.

Explicit

- Explicitly given
- Pro-actively acquired
- Expensive to collect

Implicit

- Indirectly given
- Larger quantity
- Latent qualities



Explicit or Implicit?

The screenshot shows a sidebar with filtering options. A large circle highlights the 'By Price' section, which includes categories like 'Under \$50 (18)', '\$50 to \$100 (21)', '\$100 to \$200 (14)', '\$200 to \$400 (11)', '\$400 to \$500 (12)', and a 'more' link. Below it is the 'By Rating' section, which lists star ratings from 1 to 5 stars with their respective counts: (57), (34), (17), (3), (1), and (2).

Product Rating

(1 Review)

[Read 1 Review](#)

Samsung Refrigerator Water
Regular Price \$39.99
Sale Price \$30.88

Samsung Refrigerator Water I
And French Door Refrigerator
In Stock / Free Shipping

Product Rating

(29 Review)

[Read 29 Reviews](#)

Samsung 55" Series 7 LED I
Regular Price \$3,299.99
Your Price \$2,497.00
(After \$800.00 Savings)

Explicit or Implicit?

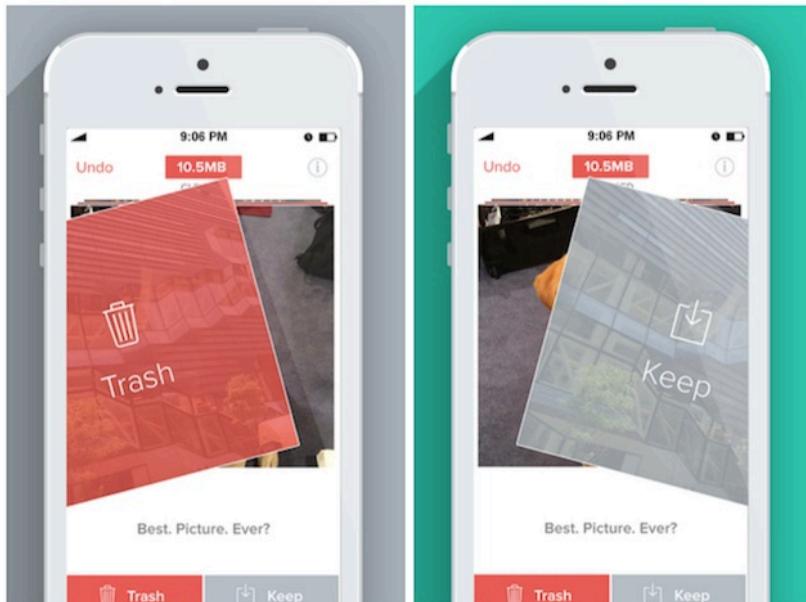
The screenshot shows a product listing page with a sidebar on the left containing filtering options. The sidebar includes sections for 'By Price' and 'By Rating'. The 'By Price' section lists price ranges with counts: Under \$50 (18), \$50 to \$100 (21), \$100 to \$200 (14), \$200 to \$400 (11), \$400 to \$500 (12), and a 'more' link. The 'By Rating' section displays a list of star rating icons with their respective counts: 5 stars (57), 4 stars (34), 3 stars (17), 2 stars (3), 1 star (1), and 0 stars (2). The main content area shows three products: 1. Samsung Refrigerator Water Filter French Door Refrigerator, Regular Price \$39.99, Sale Price \$30.88. 2. Samsung Refrigerator Water Filter French Door Refrigerator, In Stock / Free Shipping, Product Rating 4 stars (29 reviews), Read 29 Reviews. 3. Samsung 55" Series 7 LED TV, Regular Price \$3,299.99, Your Price \$2,497.00 (After \$500.00 Savings).

Explicit or Implicit?

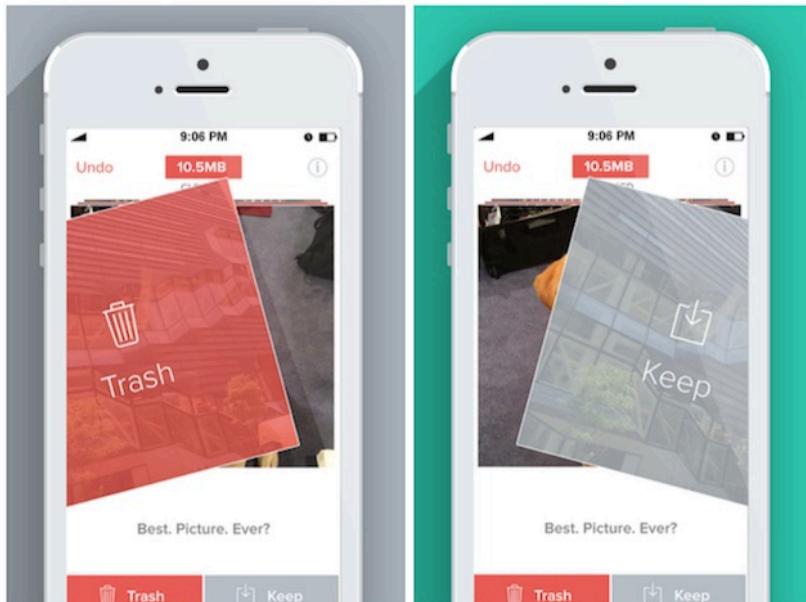
Ratings: *Explicit*

EXAMPLES – TYPES OF DATA

15

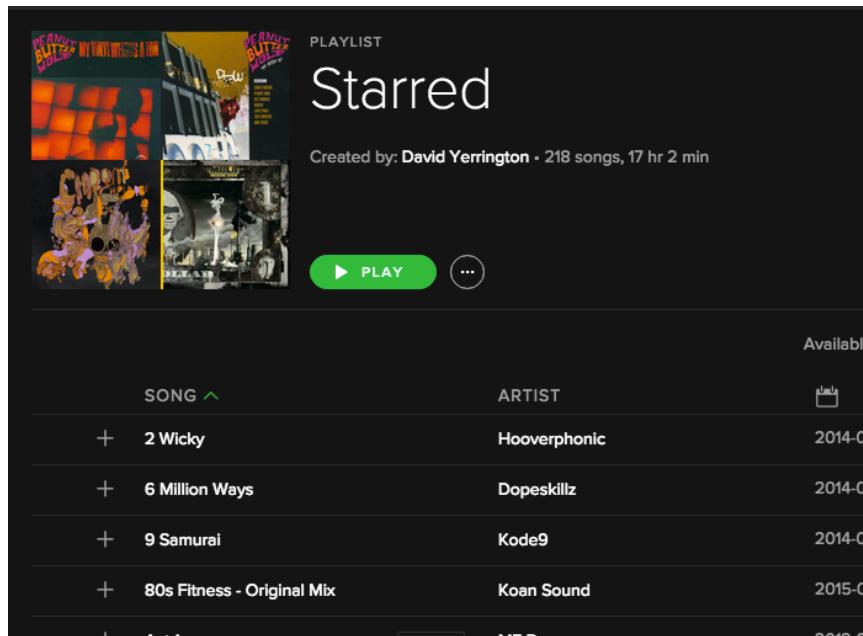


Explicit or Implicit?

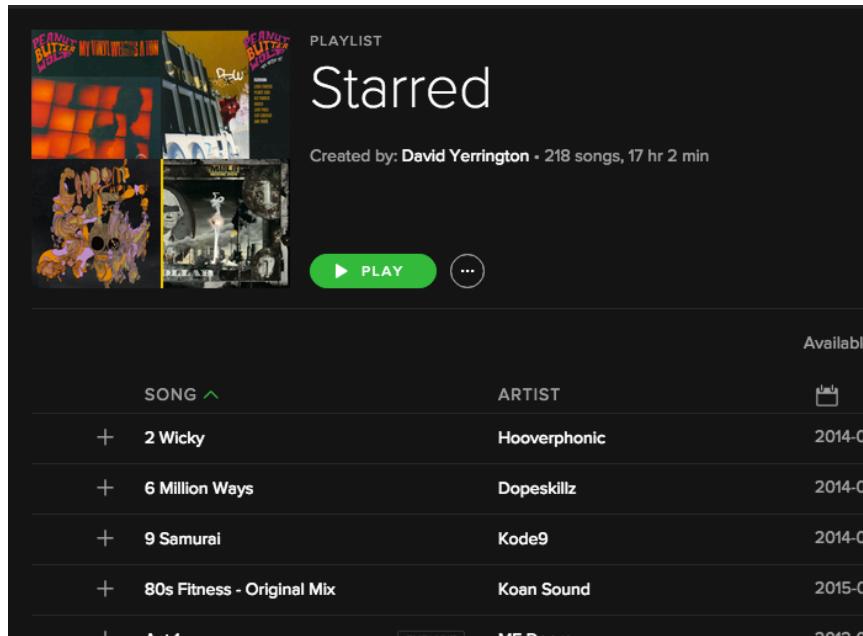


Explicit or Implicit?

Swipes: *Explicit*



Explicit or Implicit?



Explicit or Implicit?

Both!



Welcome to Target

Free Wi-Fi

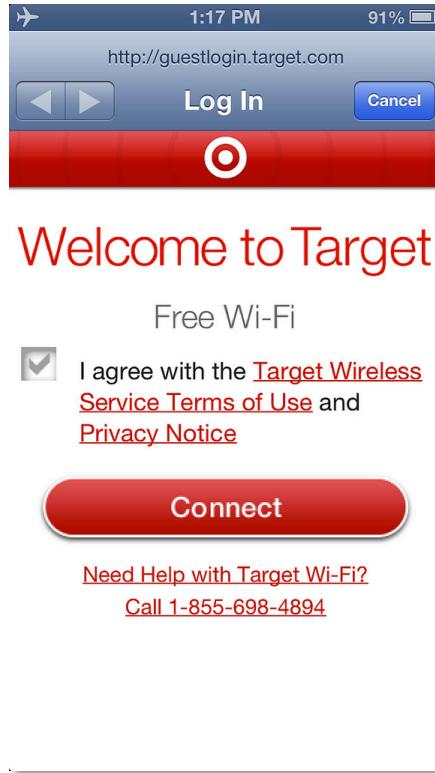


I agree with the [Target Wireless Service Terms of Use](#) and [Privacy Notice](#)

Connect

[Need Help with Target Wi-Fi?](#)
[Call 1-855-698-4894](#)

Explicit or Implicit?



Explicit or Implicit?

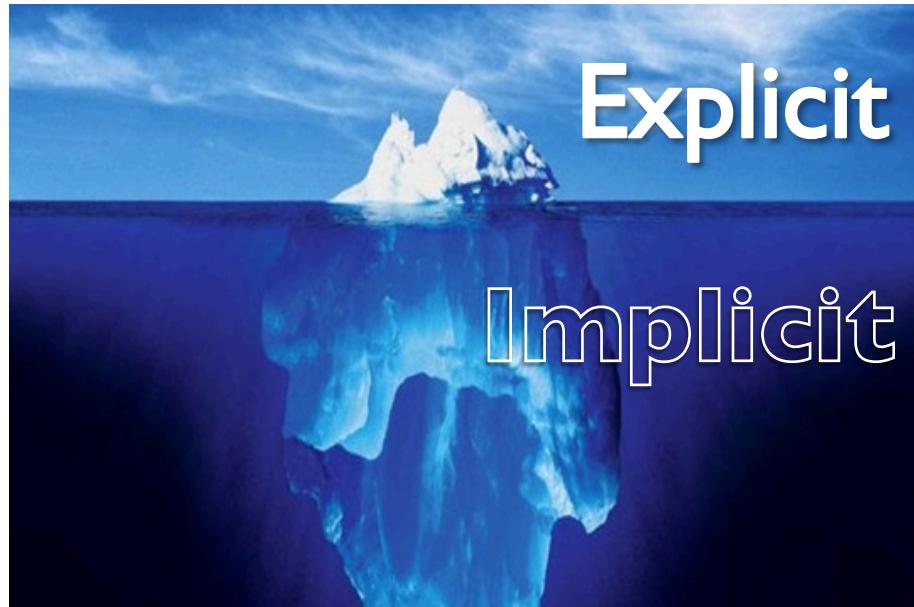
Wifi logs: *Implicit!*

Explicit

- Ie: Ratings, surveys, reviews
- Easy to interpret
- Expensive

Implicit

- Ie: Activity logs, clicks, impressions
- Hard to interpret
- Cheap



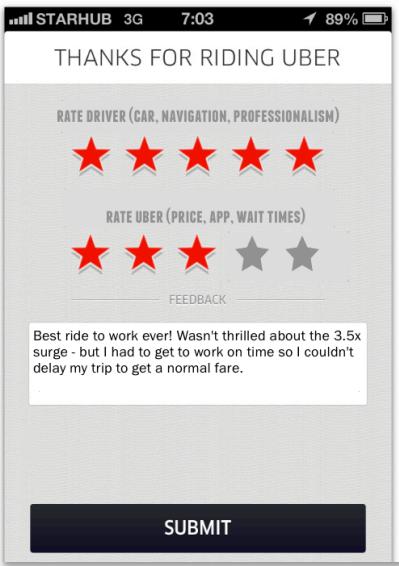
NOTE

Implicit data collection can involve some privacy issues; any system that would make recommendations must avoid overstepping its bounds.

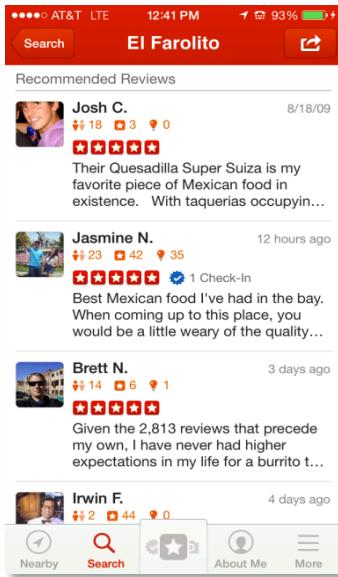
INTRO TO DATA SCIENCE

IA. EXPLICIT AND IMPLICIT FEEDBACK

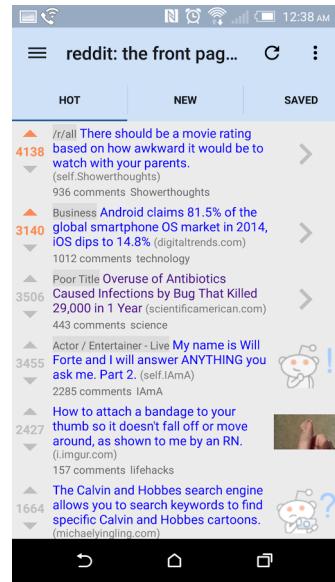
EXAMPLES – TYPES OF DATA



Uber



Yelp



Reddit

Ratings, Votes, Reviews

Detailed Seller Ratings (last 12 months)	
Criteria	Average rating
Item as described	★★★★★
Communication	★★★★★
Shipping time	★★★★★
Shipping and handling charges	★★★★★

Ebay

Explicit Feedback

- Frequently in the form of ratings
- Granularly represents preferences
- Requires extra effort from the user

Explicit Feedback Questions

- What does a rating mean?
- Do user preferences change?
- Is what is known about the data accurate?
 - Is what is collected reflect a preference at all?
 - Is it representative to the goal or only reflective of a singular characteristic?

Explicit Feedback - Considerations

- Consistent scale for all ratings
- Can ratings be skewed by self/selection-bias
- Consider the ephemeral nature of preferences
- When the data was collected
 - Before or after experience
- Context of presentation

Implicit Examples

Your Orders
Orders include Kindle book orders and any single issue purchases of newspapers and magazines.
[Learn more](#)

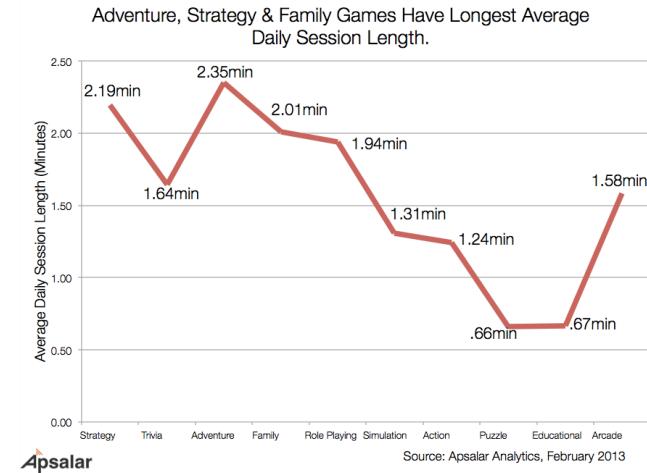
Find by author or title... View: All Books Magazines Newspapers

Title	Author	Order date
SuperFreakonomics	Steven D. Levitt, Stephen J. Dubner	December 25, 2009
The 4-Hour Workweek, Expanded and Updated: Expanded and Updated, With Over 100 New Pages of Cutting-Edge Content.	Timothy Ferriss	December 17, 2009
In Defense of Food	Michael Pollan	December 17, 2009
Kindle User's Guide, 4th Ed.	Amazon.com	December 17, 2009

« Previous | Page: 1 | Next »

A dropdown menu is open over the third row, showing delivery options:

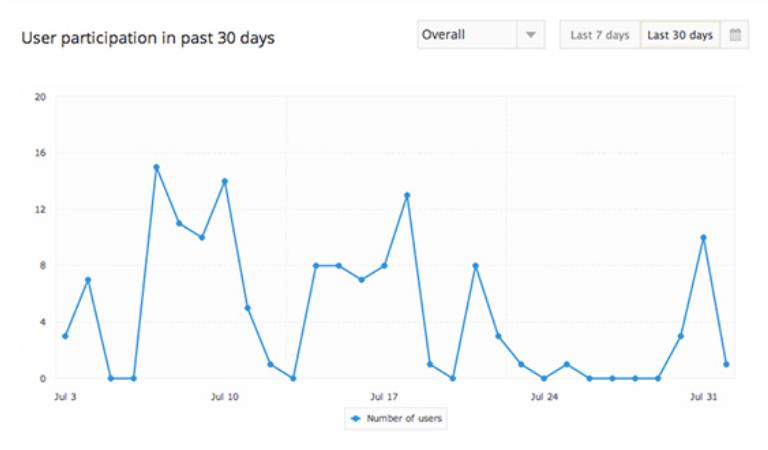
- ✓ Deliver to...
Zack's Kindle
Shapiro's Touch
-OR-
Transfer via computer...



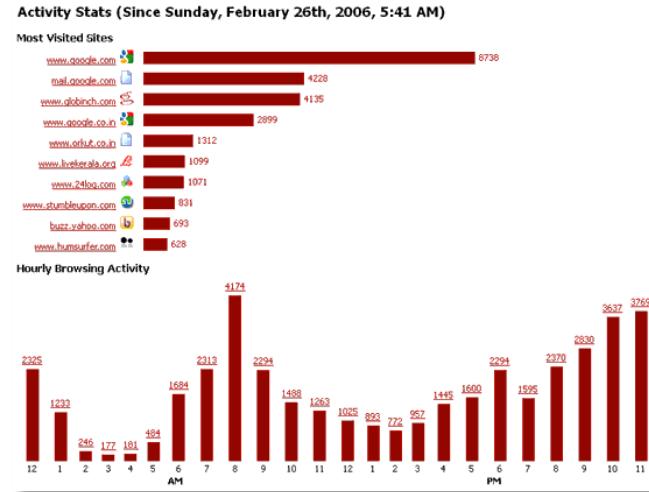
Order History

Session Length

Implicit Examples



Engagement Metrics



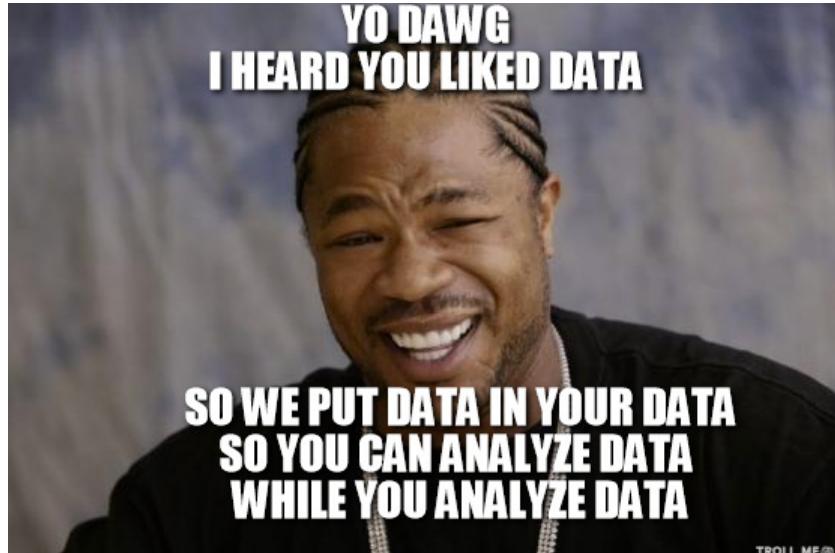
Session Length

Implicit Feedback

It's still possible to make recommendations when no rating data is explicitly collected from a user.

The goal is to convert user behavior into user preferences, but it entails one challenge: How exactly does one infer preference based on actions in a system? This can be a difficult question to answer.

Implicit Feedback



There's tons of it!

Implicit feedback is everywhere.

- Email impressions
- Email click-throughs
- Conversions
- Demographic
- Session lengths
- Login attempts
- Track plays
- Money spent
- Ad impressions
- Ad clicks
- Ad click-purchase
- Web “click depth”
- # of swipes
- Profile views
- Message initiations
- Poll Votes
- Friend / unfriend
- Follow / unfollow
- *Like
- Post text
- Image EXIF
- Friends in common
- Message text
- Food purchases
- Geospatial data
- Store cameras
- Wifi logins / MAC
- Time series
- Objects in photos
- Driving record
- Credit history
- Topics most read

Implicit feedback is valuable depending on how you look at it.



?



Implicit Feedback Caveats

Implicit Feedback Caveats

(ie: Users don't tell you what you want to know.)

- Preferences can be vague
- You may need to process tons of data to get what you want
- Analysis can be complicated / meaning hard to find
- Identities can be indistinguishable
- Users don't tell you what you want to know
- Easy to project bias onto data
- Positive / negative experience hard to assess

Implicit Feedback General Advice: Question Everything.

- Can a preference actually be observed?
- Is the lack of data actually a negative preference?
- Is there enough data to describe feedback or only a portion of it?
- Is the data scaled properly?
- Are there hidden correlations?
- Are there contradictory patterns?
- What's missing?
- Can new features be created?

Implicit + Explicit Feedback: Work together

If a user rates an item, can you use implicit feedback to validate credibility

- Did they read the article?
- Do they own the item?
- Did they rate before or after experience?
- Do other users mention them?
- Does user tend to rate high or low?
- How likely was the rating automated?

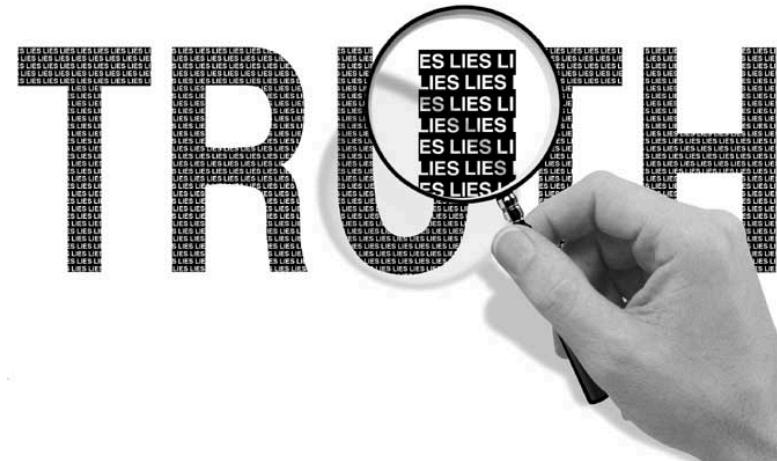
Use implicit data to understand the context and characteristics of a rating.

- Does time of day affect rating?
- Which kinds of reviews do they typically write?
- Are the reviews positive or negative?
- Do other users like their reviews?

Implicit + Explicit Feedback: Final Caveat

Take care when creating explicit data from implicit data.

- Does the set of actions reflect a preference?
- Does the scale make sense?
- Is the outcome prediction (ratings) or recommendation?



Explicit

- Higher value with respect to preferences
- Usually collected as a “rating”
- Collection is responsibility of user
- More direct evaluation of items

Implicit

- Easy to collect in large quantities
- More difficult to work with
- Assumes nothing about the user (could be anyone!)
- Goal is to convert into preferences



INTRO TO DATA SCIENCE

III. GENERAL DESIGN

There are two general approaches to the design:

There are many approaches to the design, but these are commonly modeled techniques:

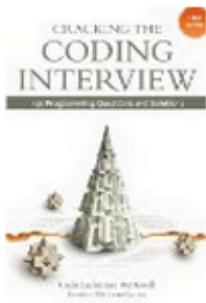
In **content-based filtering**, items are mapped into a feature space, and recommendations depend on *item characteristics*.

In contrast, an important assumption underlying all of **collaborative filtering**, is: *users who have similar preferences in the past are likely to have similar preferences in the future.*

EXAMPLES – AMAZON CONTENT-BASED

42

Recommendations for You in Books

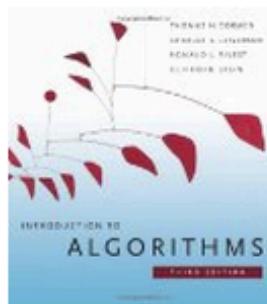


Cracking the Coding Interview: 150...
► Gayle Laakmann McDowell
Paperback

★★★★★ (166)

\$39.95 ~~\$29.95~~ \$23.22

Why recommended?

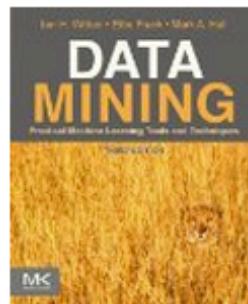


Introduction to Algorithms
Thomas H. Cormen, Charles E...
Hardcover

★★★★★ (85)

\$92.00 ~~\$80.00~~ \$80.00

Why recommended?

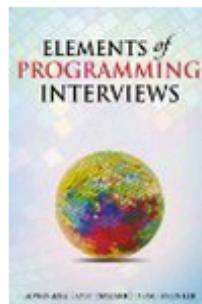


Data Mining: Practical Machine...
► Ian H. Witten, Eibe Frank, Mark A. Hall
Paperback

★★★★★ (27)

\$69.95 ~~\$42.09~~ \$42.09

Why recommended?

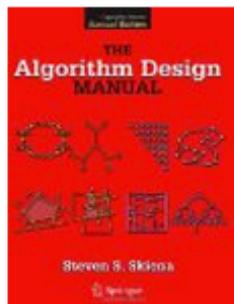


Elements of Programming Interviews...
► Amit Prakash, Adnan Aziz, Tsung-Hsien Lee
Paperback

★★★★★ (25)

\$29.99 ~~\$26.18~~ \$26.18

Why recommended?



The Algorithm Design Manual
► Steve Skiena
Paperback

★★★★★ (47)

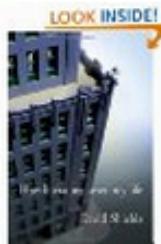
\$89.95 ~~\$71.84~~ \$71.84

Why recommended?

Customers Who Bought This Item Also Bought



Pitch Dark (NYRB Classics)
► Renata Adler
Paperback
\$11.54



How Literature Saved My Life
► David Shields
★★★★★ (60)
Hardcover
\$18.08



Bleeding Edge
Thomas Pynchon
Hardcover
\$18.05



The Flamethrowers: A Novel
► Rachel Kushner
★★★★★ (17)
Hardcover
\$15.79

EXAMPLES – NETFLIX

44

TV Shows

Your taste preferences created this row.

TV Shows.

As well as your interest in...

The image shows a section of a Netflix interface titled "TV Shows". It includes a message about taste preferences creating the row, a "TV Shows" heading, and a "As well as your interest in..." section. Below this, there are two small thumbnail images: one for "ALWAYS SUNNY IN PHILADELPHIA" featuring a cow, and another for "LAW & ORDER: SPECIAL VICTIMS UNIT" showing a group of people. To the right of this interface are two movie posters: "LOUIE" featuring Louis C.K. in a dark coat, and "Breaking Bad" featuring a yellow-tinted portrait of Walter White.

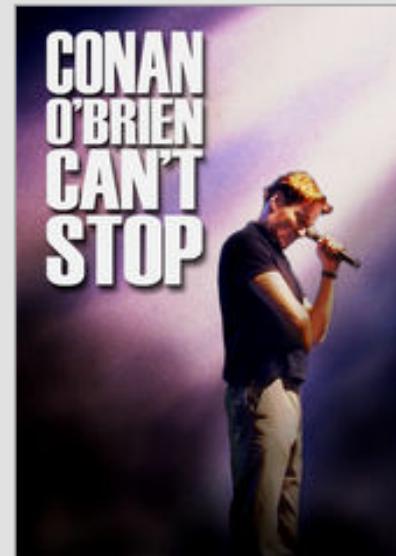
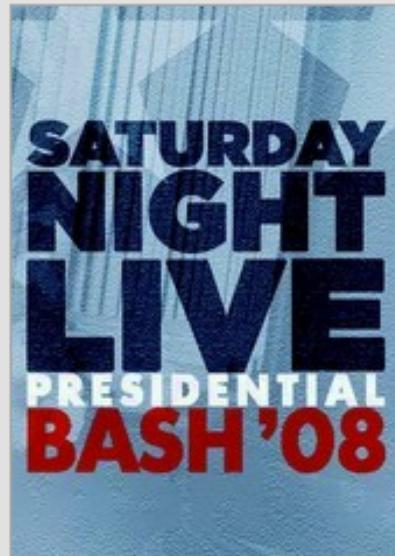
LOUIE

Breaking Bad

EXAMPLES – NETFLIX

45

Because you watched 30 Rock



EXAMPLES – YOUTUBE



Recommended for you because you watched
[Sugar Minott - Oh Mr Dc \(Studio One\)](#)



Mikey Dread - Roots and Culture
by klaxonklaxon · 1,164,133 views

Lyrics:
Now here comes a special request
To each and everyone

6:00



Recommended for you because you watched
[Thelonious Monk Quartet - Monk In Denmark](#)



Bill Evans Portrait in Jazz (Full Album)

by hansgy1 · 854,086 views

Bill Evans Portrait in Jazz 1960
1. Come Rain or Come Shine - 3.19 (0:00)
2. Autumn Leaves - 5.23 (3:24)

42:26



Recommended for you because you watched
[Bob Marley One Drop](#)



Bob Marley - She's gone

by Dionysios29 · 1,058,704 views

This is one of the eleven songs of album Kaya that Bob Marley and The Wallers creative in 1978.

Lyrics:

2:53

How can we find good recommendations?

47

- Manual Curation



- Manually Tag Attributes



content-based
filtering

- Audio Content,
Metadata, Text Analysis



- Collaborative Filtering



The image shows a screenshot of a NYTimes.com sidebar. At the top, there are two sections: 'MOST E-MAILED' on the left and 'RECOMMENDED FOR YOU' on the right. Below these sections is a list of six news items, each consisting of a number, a title in blue, and a subtitle in black.

- 1. **How Big Data Is Playing Recruiter for Specialized Workers**
- 2. SLIPSTREAM
When Your Data Wanders to Places You've Never Been
- 3. MOTHERLODE
The Play Date Gun Debate
- 4. **For Indonesian Atheists, a Community of Support Amid Constant Fear**
- 5. **Justice Breyer Has Shoulder Surgery**
- 6. BILL KELLER
Erasing History

8. How do you determine my Most Read Topics?

[Back to top ▲](#)

Each NYTimes.com article is assigned topic tags that reflect the content of the article. As you read articles, we use these tags to determine your most-read topics.

To search for additional articles on one of your most-read topics, click that topic on your personalized Recommendations page. To learn more about topic tags, visit [Times Topics](#).

NOTE

Collaborative or Content based?

8. How do you determine my Most Read Topics?

[Back to top ▲](#)

Each NYTimes.com article is assigned topic tags that reflect the content of the article. As you read articles, we use these tags to determine your most-read topics.

To search for additional articles on one of your most-read topics, click that topic on your personalized Recommendations page. To learn more about topic tags, visit [Times Topics](#).

NOTE

Collaborative or Content based?

CONTENT BASED ☺

INTRO TO DATA SCIENCE

I. CONTENT-BASED FILTERING

Content-based filtering begins by mapping each item into a feature space. Both users and items are represented by vectors in this space.

Content-based filtering begins by mapping each item into a feature space. Both users and items are represented by vectors in this space.

Item vectors measure the degree to which the item is described by each feature, and ***user vectors*** measure a user's preferences for each feature.

Content-based filtering begins by mapping each item into a feature space. Both users and items are represented by vectors in this space.

Item vectors measure the degree to which the item is described by each feature, and **user vectors** measure a user's preferences for each feature.

Ratings are generated by taking **dot products** of user & item vectors.

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

Users:

Alice = (-3, 2, -2)

Bob = (4, -3, 5)

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

Prediction (for Alice)

$$5 * -3 + 5 * 2 + 2 * -2 = -9$$

User:

Alice = (-3, 2, -2)

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

Prediction (for Alice)

$5*-3 + 5*2 + 2*-2 = -9$

$3*-3 + -5*2 + 5*-2 = -29$

User:

Alice = (-3, 2, -2)

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

Prediction (for Alice)

$5 * -3 + 5 * 2 + 2 * -2 = -9$

$3 * -3 + -5 * 2 + 5 * -2 = -29$

$-4 * -3 + -5 * 2 + -5 * -2 = +12$

User:

Alice = (-3, 2, -2)

EXAMPLE – CONTENT-BASED FILTERING

60

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

Prediction (for Alice)

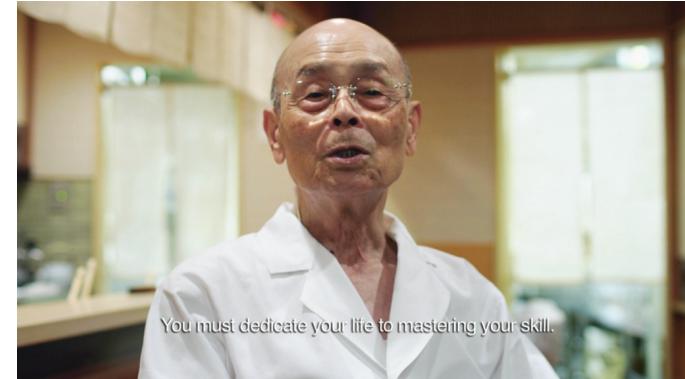
$5*-3 + 5*2 + 2*-2 = -9$

$3*-3 + -5*2 + 5*-2 = -29$

$-4*-3 + -5*2 + -5*-2 = +12$

User:

Alice = (-3, 2, -2)



You must dedicate your life to mastering your skill.

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

Prediction (for Bob)

User:

Bob = (4, -3, 5)

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

Prediction (for Bob)

$5*4 + 5*-3 + 2*5 = +15$

$3*4 + -5*-3 + 5*5 = +52$

$-4*4 + -5*-3 + -5*5 = -26$

User:

Bob = (4, -3, 5)

features = (big box office, aimed at kids, famous actors)

Items (movies):

Finding Nemo = (5, 5, 2)

Mission Impossible = (3, -5, 5)

Jiro Dreams of Sushi = (-4, -5, -5)

Prediction (for Bob)

$5*4 + 5*-3 + 2*5 = +15$

$3*4 + -5*-3 + 5*5 = +52$

$-4*4 + -5*-3 + -5*5 = -26$

User:

Bob = (4, -3, 5)

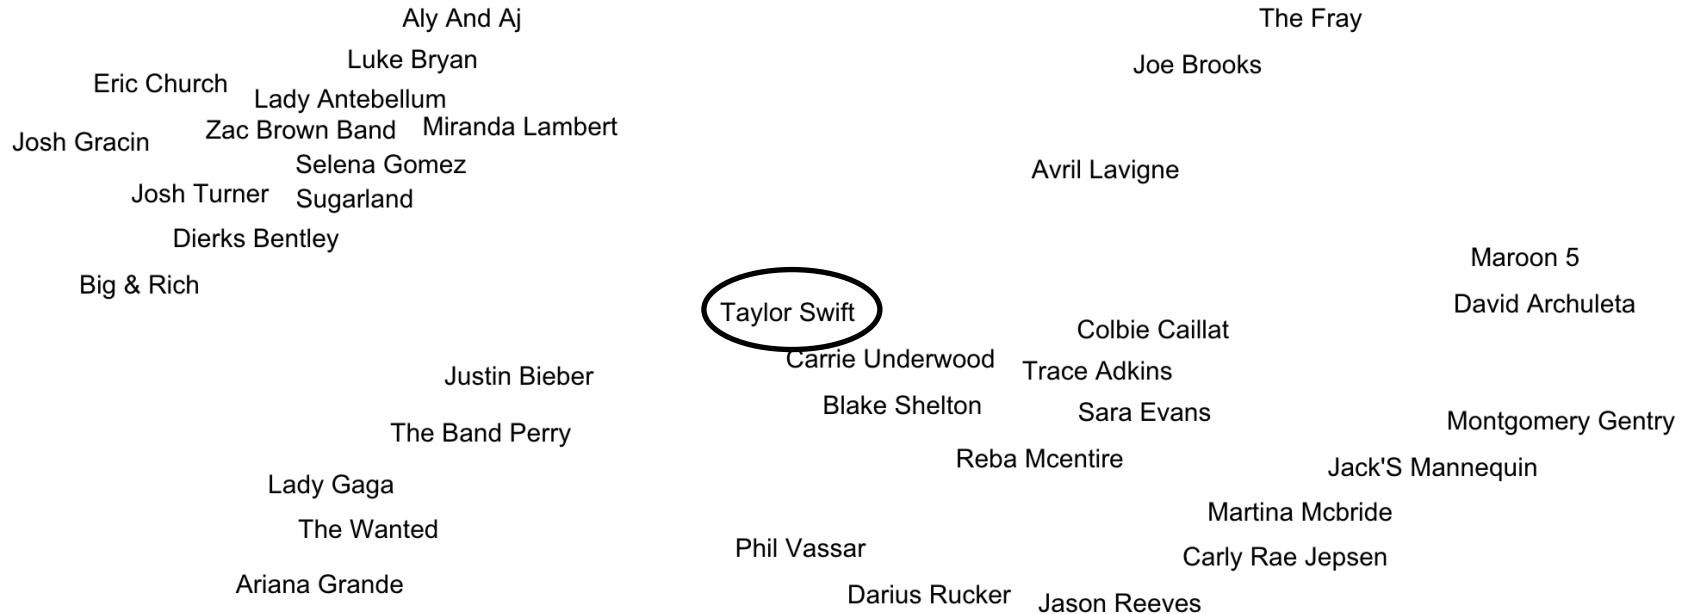


One notable example of content-based filtering is Pandora, which maps songs into a feature space using features (or “genes”) designed by the Music Genome Project.

Using song vectors that depend on these features, Pandora can create a station with music having similar properties to a song the user selects.

VISUALIZATION OF SIMILAR ARTISTS

65



THE FRAY © 2011

Content-based filtering has some difficulties:

Content-based filtering has some difficulties:

- Must map items into a feature space (usually by hand!)
- Recommendations are limited in scope (items must be similar to each other)
- Hard to create cross-content recommendations (eg books/music films...this would require comparing elements from different feature spaces!)

INTRO TO DATA SCIENCE

II. COLLABORATIVE FILTERING

Collaborative filtering refers to a family of methods for predicting ratings where instead of thinking about users and items in terms of a feature space, we are *only* interested in the existing user-item ratings themselves.

Collaborative filtering refers to a family of methods for predicting ratings where instead of thinking about users and items in terms of a feature space, we are *only* interested in the existing user-item ratings themselves.

In this case, our dataset is a *ratings matrix* whose columns correspond to items, and whose rows correspond to users.

Collaborative filtering refers to a family of methods for predicting ratings where instead of thinking about users and items in terms of a feature space, we are *only* interested in the existing user-item ratings themselves.

NOTE

The idea here is that users get value from recommendations based on other users with similar *tastes*.

In this case, our dataset is a *ratings matrix* whose columns correspond to items, and whose rows correspond to users.

RATINGS MATRIX

72

18,000 movies						
480,000 users	x	1	1	x	...	x
	x	x	x	5	...	x
	x	x	3	x	...	x
	x	4	3	x	...	2
	...	x	x	x	...	x
	x	5	x	1	...	x
	x	x	3	3	...	x
	x	1	x	x	...	2

NOTE

This matrix will always be *sparse*!

Main difference between content and collaborative filtering:

Content Based:

maps items and users into a feature space

Collaborative:

relies on previous user-item ratings

We will look at collaborative filtering in a user-user sense.

We will look at collaborative filtering in a user-user sense.

We will take a given user, and find the K most similar users, and then recommend brands from the similar users!

We will look at collaborative filtering in a user-user sense.

We will take a given user, and find the K most similar users, and then recommend brands from the similar users!

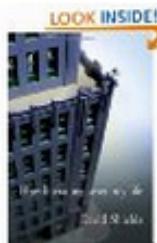
NOTE

Sound familiar? It's similar to KNN!

Customers Who Bought This Item Also Bought



Pitch Dark (NYRB Classics)
► Renata Adler
Paperback
\$11.54



How Literature Saved My Life
► David Shields
 (60)
Hardcover
\$18.08



Bleeding Edge
Thomas Pynchon
Hardcover
\$18.05



The Flamethrowers: A Novel
► Rachel Kushner

Hardcover
\$15.79

The system cannot draw inferences because it hasn't gathered enough information yet.

The cold start problem arises because we've been relying only on ratings data, or on explicit feedback from users.

The cold start problem arises because we've been relying only on ratings data, or on explicit feedback from users.

Until users rate several items, we don't know anything about their preferences!

The cold start problem arises because we've been relying only on ratings data, or on explicit feedback from users.

Until users rate several items, we don't know anything about their preferences!

We can get around this by enhancing our recommendations using implicit feedback, which may include things like item browsing behavior, search patterns, purchase history, etc.

While explicit feedback (ratings, likes, purchases) leads to high quality ratings, the data is sparse and cold starts are problematic.

While explicit feedback (ratings, likes, purchases) leads to high quality ratings, the data is sparse and cold starts are problematic.

Meanwhile implicit feedback (browsing behavior, etc.) leads to less accurate ratings, but the data is much more dense (and less invasive to collect).

INTRO TO DATA SCIENCE

III. PYTHON EXAMPLE

Our data:

- CSV of two columns, user ID and Brand
- Each row represents a user liking a brand

Example:

User ID	Brand
86509	H&M
86509	Target
86509	Old Navy
86510	H&M
86510	Lowe's
86510	Home Depot
86511	Banana Republic
86511	Kohl's
86511	Old Navy

How do we define “similarity” of users?

How do we define “similarity” of users?

This is required if we want to do user-based collaborative filtering

MATH



ALERT!!

How do we define “similarity” of users?

Jaccard Similarity:

Defines similarity between two sets of objects

How do we define “similarity” of users?

Jaccard Similarity:

Defines similarity between two sets of objects

$$JS(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

How do we define “similarity” of users?

Jaccard Similarity:

Defines similarity between two sets of objects

$$JS(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Number of similar elements

Number of distinct elements

$$JS(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

$$JS(\{1, 2, 3\}, \{2, 3, 4\}) = \{2, 3\} \quad 2$$

----- = -----

$$\{1, 2, 3, 4\} \quad 4$$

$$JS(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Exercise:

User one: {"Target", "Banana Republic", "Old Navy"}

User two: {"Banana Republic", "Gap", "Kohl's"}

JS (User one, User two) =

$$JS(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Exercise:

User one: {"Target", "Banana Republic", "Old Navy"}

User two: {"Banana Republic", "Gap", "Kohl's"}

JS (User one, User two) = 1 / 5 = .2

PYTHON ALGORITHM STEPS

1. Get list of known users in a dictionary where the key is the user ID, and the value is a list of brands they like
Example: { '83065' : ["Kohl's", 'Target'] }
2. For a given user, we will calculate their closeness to every user in csv
3. We will choose the K most similar users
4. Recommend brands liked by similar users

PYTHON ALGORITHM STEPS

1. Get list of known users in a dictionary where the key is the user ID, and the value is a list of brands they like

Example: { '83065' : ["Kohl's", 'Target'] }
2. For a given user, we will calculate their closeness to every user in csv
3. We will choose the K most similar users
4. Recommend brands liked by similar users

Consider this a kind of KNN but instead of Euclidean Distance, we are using the Jaccard Similarity

INTRO TO DATA SCIENCE

IV. THE NETFLIX PRIZE

The Netflix prize was a competition to see if anyone could make a 10% improvement to Netflix's recommendation system (accuracy measured by RMSE).

The Netflix prize was a competition to see if anyone could make a 10% improvement to Netflix's recommendation system (accuracy measured by RMSE).

The grand prize was \$1m dollars

The Netflix prize was a competition to see if anyone could make a 10% improvement to Netflix's recommendation system (accuracy measured by RMSE).

The grand prize was \$1m dollars

The ratings matrix contained >100mm numerical entries (1-5 stars) from ~500k users across ~17k movies. The data was split into train/quiz/test sets to prevent overfitting on the test data by answer submission (this was a clever idea!)

The competition began in 2006, and the grand prize was eventually awarded in 2009. The winning entry was a stacked ensemble of 100's of models (including neighborhood & matrix factorization models) that were blended using boosted decision trees.

Ultimately, the competition ended in a photo finish. The winning strategy came down to last-minute team mergers & creative blending schemes to shave 3rd & 4th decimals off RMSE (concerns that would not be important in practice).

The competition began in 2006, and the grand prize was eventually awarded in 2009. The winning entry was a stacked ensemble of 100's of models (including neighborhood & matrix factorization models) that were blended using boosted decision trees.

Ultimately, the competition ended in a photo finish. The winning strategy came down to last-minute team mergers & creative blending schemes to shave 3rd & 4th decimals off RMSE (concerns that would not be important in practice).

Though they adopted some of the modeling techniques that emerged from the competition, Netflix never actually implemented the prizewinning solution.

Why do you think that's true?