ARIMA Modeling

Matt Brems

GENERAL ASSEMBLY

Learning Objectives

- Describe the purpose of the autoregressive and moving average components.
- Define hyperparameters *p*, *d*, and *q*.
- Describe AIC.
- Find the right value of *p* and *q* using AIC.
- Find the right value of *d* using the augmented Dickey-Fuller test.
- Complete a manual GridSearch.
- Fit an ARIMA model.

We have multiple approaches to work with time series data.

Linear Models

ARIMA Models

Exponential Smoothing Methods

Recurrent Neural Networks (RNNs)

Note: This is not an exhaustive list of models, but lists the most common ones!

Why ARIMA?

- Among the most common approaches to time series modeling.
- Highly flexible; it can model time series with varying characteristics.
  - It takes information from **both long-term trends and sudden shocks!**
- Can easily be extended into more advanced models.
- Tends to perform well with moderate amounts of data.
  - It can be hard to get lots of time series data!

Downsides of ARIMA Models

- ARIMA models are best suited for **short-term forecasts**, but very quickly will start predicting the mean.
    - Some extensions to ARIMA models can handle this better.

What is an ARIMA model?

ARIMA

What is an ARIMA model?

ARIMA

Autoregressive Integrated Moving Average

What do you think the word “autoregressive” means?

What do you think the word “autoregressive” means?

- Autocorrelation is the correlation of one variable with itself.
- An autobiography is a book written by a person, about that same person.
- An autotransplant is a surgical procedure in which an organ is transplanted from a person to that same person.
- **Autoregressive means we regress a variable on itself.**
  - We’ll regress newer values on older values.

AR(p): An autoregressive model of order *p*

$$Y_t = \beta_0 + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \cdots + \beta_p Y_{t-p}$$

$$= \beta_0 + \sum_{k=1}^{p} \beta_p Y_{t-p}$$

AR(p): An autoregressive model of order *p*

$$Y_t = \beta_0 + \beta_1 Y_{t-1} + \beta_2 Y_{t-2} + \cdots + \beta_p Y_{t-p}$$

$$= \beta_0 + \sum_{k=1}^{p} \beta_p Y_{t-p}$$

Purpose: An autoregressive model explains **long-term trends** in our data.

Hyperparameter: *p*, the number of previous values of *Y* to put into our model.

We'll **GridSearch** to find this value!

What is an ARIMA model?

ARIMA

Autoregressive Integrated Moving Average

Moving Average Models

- A moving average model takes **previous error terms as inputs**.
- The goal is to predict future values based on recent forecasting errors.
  - This isn't identical to boosting, but is similar in that fitting is driven by errors.
- *Annoying: this isn't the same thing as moving average smoothing.*

MA(q): A moving average model of order *q*

$$Y_t = \mu + w_1 \varepsilon_{t-1} + w_2 \varepsilon_{t-2} + \cdots + w_p \varepsilon_{t-q}$$

$$= \mu + \sum_{k=1}^{q} w_q \varepsilon_{t-q}$$

MA(q): A moving average model of order *q*

$$Y_t = \mu + w_1 \varepsilon_{t-1} + w_2 \varepsilon_{t-2} + \cdots + w_p \varepsilon_{t-q}$$

$$= \mu + \sum_{k=1}^{q} w_q \varepsilon_{t-q}$$

Purpose: A moving average model explains **sudden shocks** in our data.

Hyperparameter: *q*, the number of previous errors ε to put into our model.

We'll **GridSearch** to find this value!

How do we GridSearch to find the best values of *p* and *q*?

- Because we're working in statsmodels, we will **manually GridSearch** values of *p* and *q* to see which gives us the **lowest AIC**.

- AIC, or Akaike Information Criterion, is a common way to evaluate time series models. (AIC is an attribute in **`statsmodels`**.)

- Remember that a model is a simplification of reality?
  - AIC attempts to measure how much information we lose when we simplify reality with a model.

$$AIC = 2 \times [\#\text{ of model parameters}] - 2 \times \log(\text{likelihood})$$

What is an ARIMA model?

ARIMA

Autoregressive Integrated Moving Average

At this point, it's helpful to see what an ARIMA model is.

$$Y_t^{(d)} = \beta_0 + \sum_{k=1}^{p} \beta_k Y_{t-k}^{(d)} + \sum_{i=1}^{q} w_i \varepsilon_{t-i} + \varepsilon_t$$

At this point, it's helpful to see what an ARIMA model is.

ARIMA

We literally just added together the AR(p) and MA(q) models.

$$Y_t^{(d)} = \beta_0 + \sum_{k=1}^{p} \beta_k Y_{t-k}^{(d)} + \sum_{i=1}^{q} w_i \varepsilon_{t-i} + \varepsilon_t$$

Autoregressive Integrated Moving Average

What is that $Y_t^{(d)}$?

- Onto the notebook!

...but first!

ARIMA Cheat Sheet

	AR	I	MA
Stands for:	Autoregressive	Integrated	Moving Average
Summary:	Regress future values on **past values**.	Differences our Y variable.	Regress future values on **past errors**.
Looks Like:	$\beta_0 + \sum_{k=1}^{p} \beta_k Y_{t-k}^{(d)}$	$Y_t^{(d)}$	$\beta_0 + \sum_{i=1}^{q} w_i \varepsilon_{t-i} + \varepsilon_t$
Purpose:	Long-term trends.	Ensure stationarity.	Sudden shocks.
Hyperparameter:	***p***	***d***	***q***
Find good value of hyperparameter by:	GridSearch	Augmented Dickey-Fuller Test	GridSearch

